

# Word2Vec Tutorial - The Skip-Gram Model

19 Apr 2016

This tutorial covers the skip gram neural network architecture for Word2Vec. My intention with this tutorial was to skip over the usual introductory and abstract insights about Word2Vec, and get into more of the details. Specifically here I'm diving into the skip gram neural network model.

## The Model

The skip-gram neural network model is actually surprisingly simple in its most basic form; I think it's the all the little tweaks and enhancements that start to clutter the explanation.

Let's start with a high-level insight about where we're going. Word2Vec uses a trick you may have seen elsewhere in machine learning. We're going to train a simple neural network with a single hidden layer to perform a certain task, but then we're not actually going to use that neural network for the task we trained it on! Instead, the goal is actually just to learn the weights of the hidden layer—we'll see that these weights are actually the "word vectors" that we're trying to learn.

Another place you may have seen this trick is in unsupervised feature learning, where you train an auto-encoder to compress an input vector in the hidden layer, and decompress it back to the original in the output layer. After training it, you strip off the output layer (the decompression step) and just use the hidden layer--it's a trick for learning good image features without having labeled training data.

## The Fake Task

So now we need to talk about this "fake" task that we're going to build the neural network to perform, and then we'll come back later to how this

indirectly gives us those word vectors that we are really after.

We're going to train the neural network to do the following. Given a specific word in the middle of a sentence (the input word), look at the words nearby and pick one at random. The network is going to tell us the probability for every word in our vocabulary of being the "nearby word" that we chose.

When I say "nearby", there is actually a "window size" parameter to the algorithm. A typical window size might be 5, meaning 5 words behind and 5 words ahead (10 in total).

The output probabilities are going to relate to how likely it is find each vocabulary word nearby our input word. For example, if you gave the trained network the input word "Soviet", the output probabilities are going to be much higher for words like "Union" and "Russia" than for unrelated words like "watermelon" and "kangaroo".

We'll train the neural network to do this by feeding it word pairs found in our training documents. The below example shows some of the training samples (word pairs) we would take from the sentence "The quick brown fox jumps over the lazy dog." I've used a small window size of 2 just for the example. The word highlighted in blue is the input word.

Source Text	Training Samples					
<table><tr><td>The</td><td>quick</td><td>brown</td></tr></table> fox jumps over the lazy dog. ➡	The	quick	brown	(the, quick) (the, brown)		
The	quick	brown				
The <table><tr><td>quick</td><td>brown</td><td>fox</td></tr></table> jumps over the lazy dog. ➡	quick	brown	fox	(quick, the) (quick, brown) (quick, fox)		
quick	brown	fox				
The quick <table><tr><td>brown</td><td>fox</td><td>jumps</td></tr></table> over the lazy dog. ➡	brown	fox	jumps	(brown, the) (brown, quick) (brown, fox) (brown, jumps)		
brown	fox	jumps				
The <table><tr><td>quick</td><td>brown</td><td>fox</td><td>jumps</td><td>over</td></tr></table> the lazy dog. ➡	quick	brown	fox	jumps	over	(fox, quick) (fox, brown) (fox, jumps) (fox, over)
quick	brown	fox	jumps	over		

The network is going to learn the statistics from the number of times each pairing shows up. So, for example, the network is probably going to get many more training samples of ("Soviet", "Union") than it is of ("Soviet", "Sasquatch").

When the training is finished, if you give it the word “Soviet” as input, then it will output a much higher probability for “Union” or “Russia” than it will for “Sasquatch”.

## Model Details

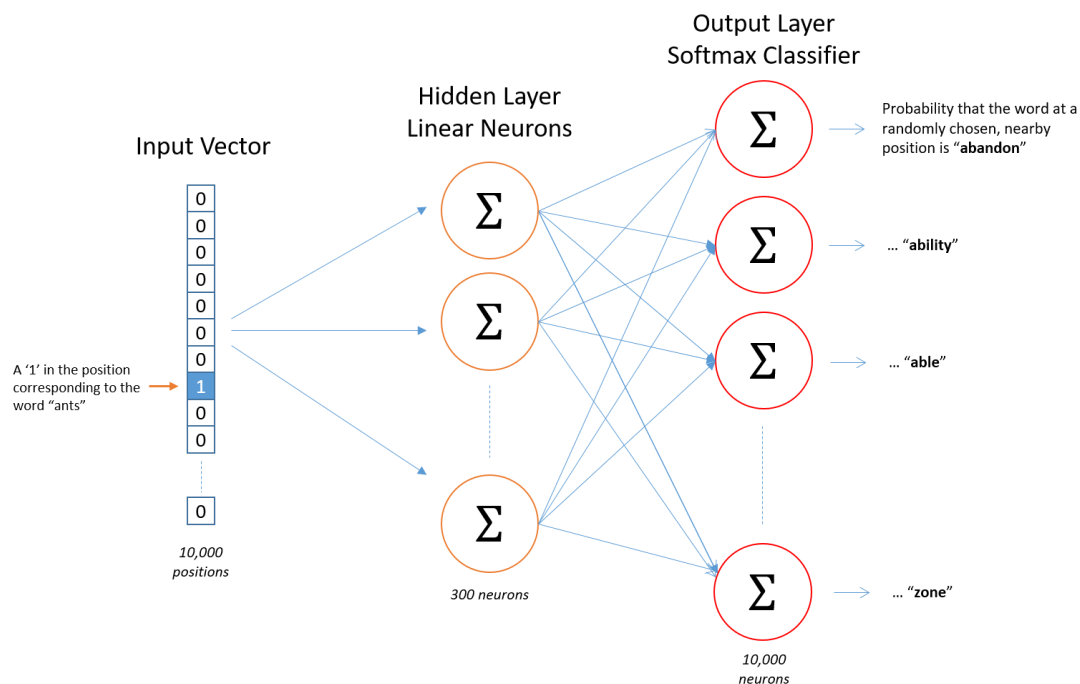
So how is this all represented?

First of all, you know you can’t feed a word just as a text string to a neural network, so we need a way to represent the words to the network. To do this, we first build a vocabulary of words from our training documents—let’s say we have a vocabulary of 10,000 unique words.

We’re going to represent an input word like “ants” as a one-hot vector. This vector will have 10,000 components (one for every word in our vocabulary) and we’ll place a “1” in the position corresponding to the word “ants”, and 0s in all of the other positions.

The output of the network is a single vector (also with 10,000 components) containing, for every word in our vocabulary, the probability that a randomly selected nearby word is that vocabulary word.

Here’s the architecture of our neural network.



There is no activation function on the hidden layer neurons, but the output neurons use softmax. We’ll come back to this later.


When *training* this network on word pairs, the input is a one-hot vector representing the input word and the training output *is also a one-hot vector* representing the output word. But when you evaluate the trained network on an input word, the output vector will actually be a probability distribution (i.e., a bunch of floating point values, *not* a one-hot vector).

## The Hidden Layer

For our example, we're going to say that we're learning word vectors with 300 features. So the hidden layer is going to be represented by a weight matrix with 10,000 rows (one for every word in our vocabulary) and 300 columns (one for every hidden neuron).

300 features is what Google used in their published model trained on the Google news dataset (you can download it from [here](#)). The number of features is a "hyper parameter" that you would just have to tune to your application (that is, try different values and see what yields the best results).

If you look at the *rows* of this weight matrix, these are actually what will be our word vectors!



So the end goal of all of this is really just to learn this hidden layer weight matrix – the output layer we’ll just toss when we’re done!

Let’s get back, though, to working through the definition of this model that we’re going to train.

Now, you might be asking yourself–“That one-hot vector is almost all zeros... what’s the effect of that?” If you multiply a  $1 \times 10,000$  one-hot vector by a  $10,000 \times 300$  matrix, it will effectively just *select* the matrix row corresponding to the “1”. Here’s a small example to give you a visual.

$$[0 \quad 0 \quad 0 \quad \mathbf{1} \quad 0] \times \begin{bmatrix} 17 & 24 & 1 \\ 23 & 5 & 7 \\ 4 & 6 & 13 \\ \mathbf{10} & \mathbf{12} & \mathbf{19} \\ 11 & 18 & 25 \end{bmatrix} = [10 \quad 12 \quad 19]$$

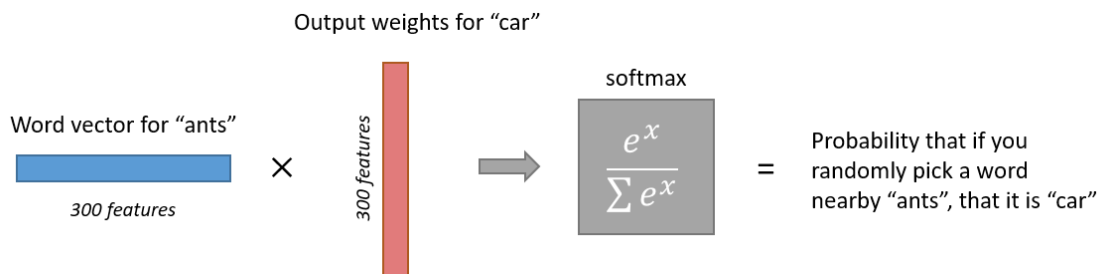
This means that the hidden layer of this model is really just operating as a lookup table. The output of the hidden layer is just the “word vector” for the input word.

# The Output Layer

The  $1 \times 300$  word vector for “ants” then gets fed to the output layer. The output layer is a softmax regression classifier. There’s an in-depth tutorial on Softmax Regression [here](#), but the gist of it is that each output neuron (one per word in our vocabulary!) will produce an output between 0 and 1, and the sum of all these output values will add up to 1.

Specifically, each output neuron has a weight vector which it multiplies against the word vector from the hidden layer, then it applies the function  $\exp(x)$  to the result. Finally, in order to get the outputs to sum up to 1, we divide this result by the sum of the results from *all* 10,000 output nodes.

Here’s an illustration of calculating the output of the output neuron for the word “car”.



Note that neural network does not know anything about the offset of the output word relative to the input word. It *does not* learn a different set of probabilities for the word before the input versus the word after. To understand the implication, let's say that in our training corpus, *every single occurrence* of the word 'York' is preceded by the word 'New'. That is, at least according to the training data, there is a 100% probability that 'New' will be in the vicinity of 'York'. However, if we take the 10 words in the vicinity of 'York' and randomly pick one of them, the probability of it being 'New' is *not* 100%; you may have picked one of the other words in the vicinity.

## Intuition

Ok, are you ready for an exciting bit of insight into this network?

If two different words have very similar “contexts” (that is, what words are likely to appear around them), then our model needs to output very similar results for these two words. And one way for the network to output similar context

predictions for these two words is if *the word vectors are similar*. So, if two words have similar contexts, then our network is motivated to learn similar word vectors for these two words! Ta da!

And what does it mean for two words to have similar contexts? I think you could expect that synonyms like "intelligent" and "smart" would have very similar contexts. Or that words that are related, like "engine" and "transmission", would probably have similar contexts as well.

This can also handle stemming for you – the network will likely learn similar word vectors for the words "ant" and "ants" because these should have similar contexts.

## Next Up

You may have noticed that the skip-gram neural network contains a huge number of weights... For our example with 300 features and a vocab of 10,000 words, that's 3M weights in the hidden layer and output layer each! Training this on a large dataset would be prohibitive, so the word2vec authors introduced a number of tweaks to make training feasible. These are covered in [part 2 of this tutorial](#).

## Other Resources

I've also created a [post](#) with links to and descriptions of other word2vec tutorials, papers, and implementations.

### Cite

McCormick, C. (2016, April 19). *Word2Vec Tutorial - The Skip-Gram Model*. Retrieved from <http://www.mccormickml.com>

158 Comments

mccormickml.com

1 Login ▾

♥ Recommend 72

🔗 Share

Sort by Best ▾



Join the discussion...

LOG IN WITH



OR SIGN UP WITH DISQUS (?)

Name



**zul waker** • 7 months ago

You are amazing. I tried to understand this from so many sources but you gave the best explanation possible. many thanks.

18 ^ | v • Reply • Share ›



**Chris McCormick** Mod ➔ zul waker • 7 months ago

Thanks so much, really glad it helped!

1 ^ | v • Reply • Share ›



**Albert Wang** • 6 months ago

The best word2vec tutorial I have ever read besides the paper.

One question:

Since the algorithm knows nothing about the slicing window, does that mean there is no difference between the first word after the target word and the second word after the target word?

For example, if the window is [I am a software engineer], here the target word is "a".

The algorithm will train the neural network 4 times. Each time, the output will be a softmax vector and it computes the cross entropy loss between the output vector and the true one-hot vector which represents "i", "am", "software", and "engineer".

Therefore, this is just a normal softmax classifier. But word2vec uses it in a smart way.

Do they use "cross entropy"? Which loss function do they use?

4 ^ | v • Reply • Share ›



**Chris McCormick** Mod ➔ Albert Wang • 6 months ago

Hi Albert,

You're correct that the position of the word within the context window has no impact on the training.

I'm hesitant to answer your question about the cost function because I'm not familiar with variations of the softmax classifier, but I believe you're correct that it's an ordinary softmax classifier like [here] (<http://ufldl.stanford.edu/t...>

To reduce the compute load they do modify the cost function a bit with something called Negative Sampling--read about that in part 2 of this tutorial.

^ | v • Reply • Share ›



**Albert Wang** ➔ Chris McCormick • 6 months ago





Thank you for replying.

I am aware of negative sampling they used. It's more like an engineering hack to speed up stuff.

They also used noise contrastive estimation as another loss function candidate.

But, I want to double confirm that ordinary softmax with full cross entropy is perfect valid in terms of computation correctness instead of efficiency.

^ | v • Reply • Share ›



**raj1514** • 8 months ago

Thanks for this post! It really saved time in going through papers about this...

3 ^ | v • Reply • Share ›



**Chris McCormick** Mod → raj1514 • 8 months ago

Great! Glad it helped.

^ | v • Reply • Share ›



**Arish Ali** • 5 months ago

Loved the simplicity of the article and the visualizations of different numeric operations made it so easy to understand

2 ^ | v • Reply • Share ›



**Chris McCormick** Mod → Arish Ali • 5 months ago

Awesome, thanks!

^ | v • Reply • Share ›



**ningyuwhut** • 3 days ago

I have a question that how to understand skip in the name "the Skip-Gram Model" literally? I mean why this model called the skip-gram model. Thanks

1 ^ | v • Reply • Share ›



**Ajay Prasadh** • 6 months ago

Explicitly defining the fake task helps a lot in understanding it. Thanks for an awesome article !

1 ^ | v • Reply • Share ›



**Chris McCormick** Mod → Ajay Prasadh • 6 months ago

Glad it helped, thanks!

^ | v • Reply • Share ›



**Nazar Dikhil** → Chris McCormick • 4 months ago



Please

I want to do predicted a fuzzy time series (fuzzification by FC-mean)

By RBF neural network

But I'm having a problem with training, can you help me ???

Thanks

^ | v • Reply • Share ›



**Ahmed EIFki** • 7 months ago

Thank you for this article, however in the hidden layer part how did you choose the number of features for the hidden layer weight matrix ?

1 ^ | v • Reply • Share ›



**Chris McCormick** Mod → Ahmed EIFki • 7 months ago

Hi, Ahmed - 300 features is what Google used in training their model on the Google news dataset. The number of features is a "hyper parameter" that you would just have to tune to your application (that is, try different values and see what yields the best results).

^ | v • Reply • Share ›



**Ahmed EIFki** → Chris McCormick • 7 months ago

Hi Chris, first thank you for answering my question and second i have another question which is related to the word representation. After initializing the Word2Vec method (in python) with the its corresponding parameters (such as the array of sentences extracted from documents, number of hidden layers, window, etc) then each word will be described by a set of values (negative and positive) in the range -1, 1 which is a vector depending on the numbers of features chosen previously. As i understood from gensim documentation all of the words representation can be extracted by means of the method `save_word2vec_format(dest_file_name, binary=False)` that can be fed later to another network. So can you confirm if what i understood if right or wrong ?

^ | v • Reply • Share ›



**Chris McCormick** Mod → Ahmed EIFki • 6 months ago

I believe that's all correct! In gensim, you can look up the word vector for a word using, e.g., .

```
model = gensim.models.Word2Vec.load_word2vec_format
model['hello']
Out[12]:
array([-0.05419922,  0.01708984, -0.00527954,  0.33
        -0.01397705, -0.15039062, -0.265625,  0.01
```

```
np.min(model['hello'])  
Out[15]: -0.59375
```

```
np.max(model['hello'])  
Out[16]: 0.50390625
```

^ | v • Reply • Share ›



**Ahmed EIFki** → Chris McCormick • 6 months ago

Hello Chris, thank you for having devoted part of your time to read and answer my question because it allows me to confirm my doubts on this API (Word2Vec). Your answers as well as your article help me a lot to progress in my project. :)

^ | v • Reply • Share ›



**Chris McCormick** Mod → Ahmed EIFki • 6 months ago

Awesome, glad I could help!

^ | v • Reply • Share ›



**Bob** • 9 months ago

Nice article, very helpful, and waiting for your negative sample article.

My two cents, to help avoid potential confusion :

First, the CODE : <https://github.com/tensorfl...>

Note though word2vec looks like a THREE-layer (i.e., input, hidden, output) neural network, some implementation actually takes a form of kind of TWO-layer (i.e., hidden, output) neural network.

To illustrate:

A THREE layer network means :

input  $\times$  matrix  $W_1$   $\rightarrow$  activation(hidden, embedding)  $\rightarrow$  times matrix  $W_2$   $\rightarrow$  softmax  $\rightarrow$  Loss

A TWO layer network means :

activation(hidden, embedding)  $\rightarrow$  times matrix  $W_2$   $\rightarrow$  softmax  $\rightarrow$  Loss

How ? In the above code, they did not use Activation( matrix  $W_1$   $\times$  input) to generate a word embedding.

Instead, they simply use a random vector generator to generate a 300-by-1 vector and use it to represent a word. They generate 5M such vectors to represent 5M words as their embeddings, say their dictionary consists of 5M words.

in the training process, not just the  $W_2$  matrix weights are updated, but also

"the EMBEDDINGS ARE UPDATED" in the back-propagation training process as well.

In this way, they trained a network where there is no matrix  $W_1$  that need to be updated in the training process.

It confused me a little bit at my first look at their code when I was trying to

it confused me a little bit at my first look at their code, when I was trying to find "two" matrices.

Sorry I had to use Capital letter as highlight to save reader's time. No offence.

1 ^ | v • Reply • Share ›



**Chris McCormick** Mod → Bob • 9 months ago

I could be wrong, but let me explain what I think you are seeing.

As I understand it, your diagram of a "3-layer network" is incorrect because it contains three weight matrices, which you've labeled  $W_1$ , word embeddings, and  $W_2$ . The correct model only contains two weight matrices--the word embeddings and the output weights.

Where I could see the *code* being confusing is in the input layer. In the mathematical formulation, the input vector is this giant one-hot vector with all zeros except at the position of the input word, and then this is multiplied against the word embeddings matrix. However, as I explained in the post, the effect of this multiplication step is simply to select the word vector for the input word. So in the actual code, it would be silly to actually generate this one-hot vector and multiply it against the word embeddings matrix--instead, you would just select the appropriate row of the embeddings matrix.

Hope that helps!

1 ^ | v • Reply • Share ›



**Chris McCormick** Mod → Bob • 8 months ago

FYI, I've written a [part 2](#) covering negative sampling.

^ | v • Reply • Share ›



**Calvin Ku** • a year ago

Thanks for the article Chris! I was going through a TensorFlow tutorial on Word2Vec and really couldn't make heads or tails of it. This article really helps a lot!

I have one question regarding the labels though. In the first figure, my understanding is, for each word (one-hot encoded vector) in the input, the NN outputs a vector of the same dimension (in this case,  $\text{dim} = 10,000$ ) in which each index contains the probability of the word of that index appearing near the input word. And since this is a supervised learning, we should have readied the labels generated from our training text, right (we already know all the probabilities from training set)? This means the labels are a vector of probabilities, and not a word, which doesn't seem to be agreed by your answer to [@Mostaphe](#).

Also I don't think the probabilities in the output vector should sum up to

one. Because we have a window of size 10 and in the extreme case, say we have a text of repeating the same sentence of three words over and over, then all the words will appear in the vicinity of any other word and they should always have probability of 1 in any case. Does this make sense?

1 ^ | v • Reply • Share ›



**Chris McCormick** Mod → Calvin Ku • a year ago

Hi Calvin, thanks, glad it was helpful!

The outputs of the Softmax layer are guaranteed to sum to one because of the equation for the output values--each output value is divided by the sum of all output values. That is, the output layer is normalized.

I get what you are saying, though, and it's a good point--I believe the problem is in my explanation.

Here is, I think, the more technically correct explanation: Let's say you take all the words within the window around the input word, and then pick one of them at random. The output values represent, for each word, the probability that the word you picked is that word.

Here's an example. Let's say in our training corpus, *every occurrence* of the word 'York' is preceded by the word 'New'. That is, at least according to the training data, there is a 100% probability that 'New' will be in the vicinity of 'York'. However, if we take the words in the vicinity of 'York' and randomly pick one of them, the probability of it being 'New' *is not* 100%.

I will add a note to my explanation; thanks for catching this!

^ | v • Reply • Share ›



**Lifemakers Studio** → Chris McCormick • 6 months ago

So, for each input word, the ideal output of trained network should be a vector of 10,000 floating-point values, all of which should be 0 except those whose words have been ever found nearby the input word, and each such non-zero value should be proportional (pre-softmax) to the number of occurrences of that word near the input word?

If so, how can this ideal training be achieved if the network looks at only one nearby word at a time? For example if "York" has been seen 100 times near "new" and 1 time near "kangaroo", and we give the network the York/kangaroo pair, wouldn't the training algorithm hike the output for "kangaroo" all the way to 1 at this step, instead of the 1/100 as it should be? Or does the fact that we'll feed it York/new pairs 100

times as often take care of this?

^ | v • Reply • Share ›



**Chris McCormick** Mod → Lifemakers Studio

• 6 months ago

Your very last statement is correct. Each training sample is going to tweak the weights a little bit to more accurately match the output suggested by that sample. There will be many more samples of the "york" and "new" combination, so that pairing will get to tweak the weights more times, resulting in a higher output value for 'new'.

^ | v • Reply • Share ›



**Emanuel** • 13 days ago

Thanks for article! I just have a question: when we calculate probability that if we pick random word in radius of word and we choose car... I understand that  $1 \times 300$  vector for ant is from hidden layer weight matrix, but is  $300 \times 1$  output vector for car also from there (but is transposed)? I believe you mean that this  $300 \times 1$  vector is vector that is fed in softmax layer, so second sigma in second picture? But how do I go, what is difference between activation of hidden layer (first sigma) and input to output layer (second sigma)? So is input to second sigma just weighted sum of hidden layer activations multiplied by weights (which we learn)? But this isn't hidden layer weight matrix (also notated look up table), what is intuitive difference in interpretation between these two, does this second weight matrix also play any important role?

Sorry if it's too much questions, but I didn't find answer in article. I appreciate any help :)

^ | v • Reply • Share ›



**Sepehr Jng** • 20 days ago

Thank you very much. This was definitely one of the best, if not the best, tutorial I've seen on this subject. However, I don't really get one part of the model. I'd really appreciate it if anyone could help:

Say we have a window of size 5 (1 center word + 2 words before and 2 words after). So Skip-gram model actually has 5 softmax units each having the size of vocabulary. also I've seen in several articles people saying that the matrix  $U$  for all of these units is the same so my confusion is that all of the inputs would be the same since we multiply  $V_c$  and the same  $U$  for all softmax units. (Meaning that the output of all the softmax units are always equal to other units) So why does it make sense to have more than one?

Please help. I really need to understand this

Thanks :)

^ | v • Reply • Share ›



**saria goudarzvand** • 22 days ago

Awesome, I really like your detailed explanations.

^ | v • Reply • Share ›



**Maryam Mohamed** • a month ago

The simplest and the best explanation. Thank you

^ | v • Reply • Share ›



**Allen Z Young** • 2 months ago

Awesome!

^ | v • Reply • Share ›



**maryem** • 2 months ago

Thank you for your clear tutorial, if you allow i have a question :  
What is the method used to calculate the weights inserted into the weight matrix?

^ | v • Reply • Share ›



**Chris McCormick** Mod → maryem • 2 months ago

The weights are learned through the neural network training procedure. Look around for a tutorial on "back propagation" or neural network training to learn more. I took Andrew Ng's Machine Learning class on Coursera to learn back propagation, but you might be able to find a simpler explanation if you're not interested in all of the details.

^ | v • Reply • Share ›



**Pi lou** • 2 months ago

Hi ! Can someone explain me what "features" are for a word ?

^ | v • Reply • Share ›



**Chris McCormick** Mod → Pi lou • 2 months ago

A word vector is simply an array of floating point values. You can refer to the values in the vector as "features". It's difficult to provide an interpretation of the feature values--I think the best you can say is that, taken all together as a word vector, words which have similar meanings will also have similar word vectors.

^ | v • Reply • Share ›



**Pi lou** → Chris McCormick • 2 months ago

Ok !!

Thank you for answering !!

^ | v • Reply • Share ›



**Zach Moshe** • 3 months ago

A great explanation! Thanks! I needed a quick reminder on word2vec internals and finally found something readable that explains exactly what I need. Going to part 2 now...

^ | v • Reply • Share ›



**Chris McCormick** Mod → Zach Moshe • 2 months ago

Thanks for the comment, glad it helped!

^ | v • Reply • Share ›



**Subhash Bylaiah** • 3 months ago

Thanks Chris.. This is an awesome tutorial.. glad I found it.. I also went through many of the comment discussions below.. thanks for taking the time to answer each of them diligently.

^ | v • Reply • Share ›



**Chris McCormick** Mod → Subhash Bylaiah • 3 months ago

Thanks, Subhash, glad it's been helpful!

^ | v • Reply • Share ›



**Aazim Lakhani** • 3 months ago

We have a matrix (10000 \* 300), where 10000 represents words & 300 represents the neurons. So, we're bound to have some rows that are identical. However, from what I understood, for a given word vector (a row) only one element of the matrix would be 1. Is it possible for 2 word vectors to have an identical row. Could you please clarify / correct my understanding.

^ | v • Reply • Share ›



**Chris McCormick** Mod → Aazim Lakhani • 3 months ago

Hi Aazim - The 10,000 words are the 10,000 words in the \*vocabulary\*--that is, they represent every \*unique\* word found in the training corpus. So, no, there shouldn't be any duplicates. Does that clarify it?

^ | v • Reply • Share ›



**Octavian Tuchila** • 3 months ago

What is the connection between the Hidden Layer Weight Matrix and the Word Vector Lookup Table?

Are they the same thing?

Or are they multiplied with each other?

Great article, by the way.

^ | v • Reply • Share ›



**Chris McCormick** Mod → Octavian Tuchila • 3 months ago





Hi Octavian, thanks, appreciate it!

Yes, they are the same thing. As a weight matrix, you might view them as 300 neurons with 10,000 weights each. But when you use it as a word vector lookup table, you view it as 10,000 word vectors with 300 features each.

1 ^ | v • Reply • Share ›



**Laurie Saint Crieg** • 3 months ago

Hello !! Very good explanation !! It's clearer now.

Is it possible to consider only word left or right?

Thank you :)

^ | v • Reply • Share ›



**Laurie Saint Crieg** → Laurie Saint Crieg • 3 months ago

And another question, is it possible to look for the neighbours of a word who is only in the training document in a new document?  
Or, are we obliged to make a learning on the both documents in the same time?

^ | v • Reply • Share ›



**Chris McCormick** Mod → Laurie Saint Crieg • 3 months ago

Once you've learned the word vectors, they can be used on documents that weren't part of the training set, yes. The only concern would be if the word is used in a significantly different way between the training versus production

---

## Related posts

[Concept Search on Wikipedia](#) 22 Feb 2017

[Getting Started with mlpack](#) 01 Feb 2017

[Word2Vec Tutorial Part 2 - Negative Sampling](#) 11 Jan 2017

---



